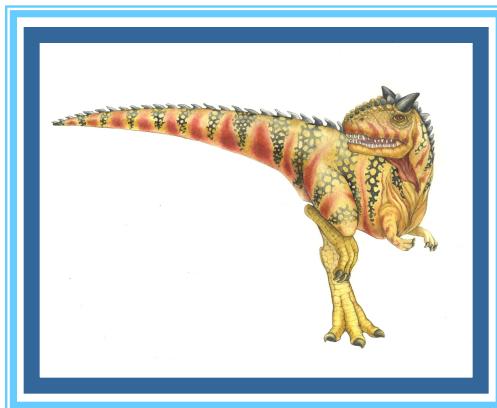


Chapter 1: Introduction (to Operating Systems)





Chapter 1: Introduction

- What we are going to discuss in this chapter:
 - What Operating Systems are & what they do.
 - Computer-System Organization & Architecture (**very brief background**)
 - Operating-System Structure (**parts of an OS**).
 - Operating-System Operations (**tasks performed by each**).
 - ▶ Process Management
 - ▶ Memory Management
 - ▶ Storage Management
 - ▶ Protection and Security
 - Some Computing Environments (traditional, mobile etc.).
 - Open-Source Operating Systems.





Objectives of Chapter 1

- To describe the basic organization of computer systems and the role of interrupts in a computer system.
- To give a tour of the major components of operating systems.
 - An introductory discussion.
 - More details about each component in the later chapters.
- Illustrate the dual-mode operation of an OS.
- To discuss how an OS is used in various computing environments.
- To discuss free and open-source operating systems.





Some Famous Operating Systems

■ For desktops & laptops?

- Microsoft Windows – 95, 98, 2000, XP, Vista, 7, 8, 10.
- Ubuntu Linux – Trusty Tahr (kind of goat), Xenial Xerus (squirrel).
- Apple Mac OS – Cheetah, Puma, Jaguar, Tiger, Leopard, Snow Leopard, Yosemite, El Capitan, Sierra, High Sierra.

■ For Servers.

- Windows Server, Red Hat Enterprise Linux Server.

■ Phones & tablets?

- Android – Jelly bean, Kit Kat, Marshmallow, Nougat.
- iOS – 8, 9, 10, 11, 12...?.

■ Watches.

- Android wear.
- iWatch OS.

■ Key point: operating systems are of various kinds, offering varied functionality (desktop vs mobile OS vs watch OS?).





What is an Operating System?

- A program that acts as an intermediary/mediator between a user of a computer and the computer hardware.
- A large piece of software (code) that sits between the hardware & the user.
- Operating system goals:
 - Execute user programs (on user's behalf).
 - Make the computer system easy to use (for users).
 - Use the computer hardware in an efficient manner.
- Note that some of the goals may be conflicting. Example?
- Ease of use may come with a side effect of less than optimal resource utilization.
- Consider a command line vs. graphic user interface (GUI).
- GUI is more user friendly, but less efficient, as some of the hardware resources are being used to support it.
- Command line, though not so easy to use, is more efficient.





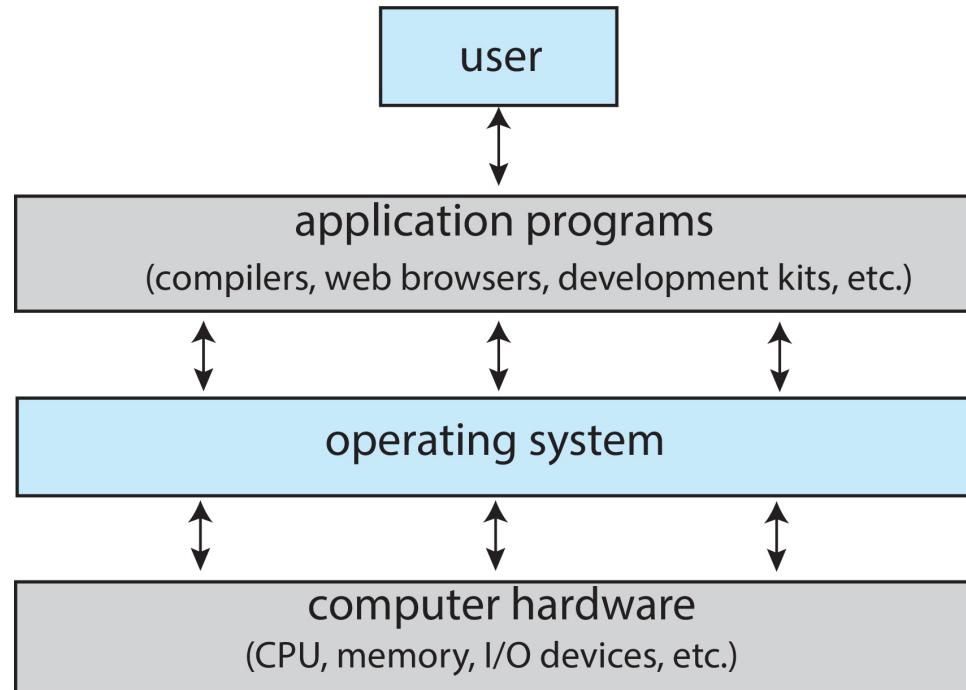
Computer System Structure

- **What's the role of an OS in the overall computer system?**
- A computer system can be divided into four components:
 - **Hardware** – provides basic computing resources.
 - ▶ CPU, memory, disk, I/O devices.
 - **Operating system.**
 - ▶ Controls and coordinates use of hardware among various applications and users
 - **Application programs** – define the ways in which the system resources are used to solve the computing problems of the users.
 - ▶ Word processors, web browsers, database systems, video games.
 - **Users**
 - ▶ People.
- Next slide shows the structure of a computer system.





Abstract View of Components of Computer





OS goals?

- From the user's view.
- Regular/common users want **convenience/ease of use** and **good performance**.
 - Don't care much about **resource utilization** (how hardware is shared, typically between multiple users).
- But, shared computers such as **mainframes/servers** must keep all users happy, it's more about **fairness** among users.
- Users of dedicated systems such as **workstations** have dedicated resources but frequently use shared resources from **servers (a combination of the above two options)**.
- Handheld computers are resource poor (compared to desktops), optimized for **usability** and **battery life**.
 - Need user interfaces such as touch screens, voice recog.
- Some computers have little or no user interface (to interact with hardware), such as embedded computers in home devices (fridge) and automobiles.





System View

■ OS is a **resource allocator**.

- Manages all hardware resources of a computer.
- Decides between conflicting requests for efficient and fair resource use.
- For example, if multiple programs want the CPU, how to allocate?

■ OS is a **control program**.

- Controls execution of programs & usage of various I/O devices to prevent errors and improper use of the computer.
- One program cannot interfere in the execution of another program. What is wrong with this interference?





What is an Operating System?



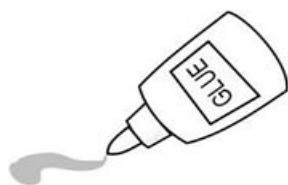
- **Referee**

- Manage sharing of resources, Protection, Isolation
 - » Resource allocation, isolation, communication



- **Illusionist**

- Provide clean, easy to use abstractions of physical resources
 - » Infinite memory, dedicated machine
 - » Higher level objects: files, users, messages
 - » Masking limitations, virtualization



Glue

- Common services
 - » Storage, Window system, Networking
 - » Sharing, Authorization
 - » Look and feel





What does an OS contain?

- Debatable, no universally accepted theory.
- “*Everything a vendor ships when you order an operating system*” – a naïve/ simplistic approximation.
 - This varies a lot.
 - Some operating systems are < 1 MB (embedded devices, a microwave oven) with little/no GUI, others are several GBs (Windows 10 needs 15 GB, iOS 12 takes up 2.25 GB).
- Another theory: “*the one program running at all times on the computer*”, called the **kernel/core**.
- Every other software is either:
 - A **system program**: ships with the operating system, but not part of the kernel, e.g. compiler, assembler.
 - An **application program**: all programs not associated with the operating system, e.g. Facebook, Twitter.

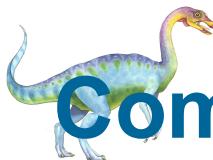




Computer Startup

- **Bootstrap program** is loaded at power-up or reboot.
 - To boot means to start.
 - Computer needs an initial program to run when it is powered on.
 - This program is typically stored in ROM (read only memory), generally known as **firmware**.
 - It typically initializes the computer system – CPU registers, device controllers, memory contents etc..
 - Bootstrap locates (on disk) & loads operating system kernel (into memory).
 - Once kernel is up & running, it waits for some user event to occur.
 - ▶ Connecting to a Wi-Fi network.
 - ▶ Opening a web browser.
 - Bootloader for Linux: GRUB (Grand Unified Bootloader).





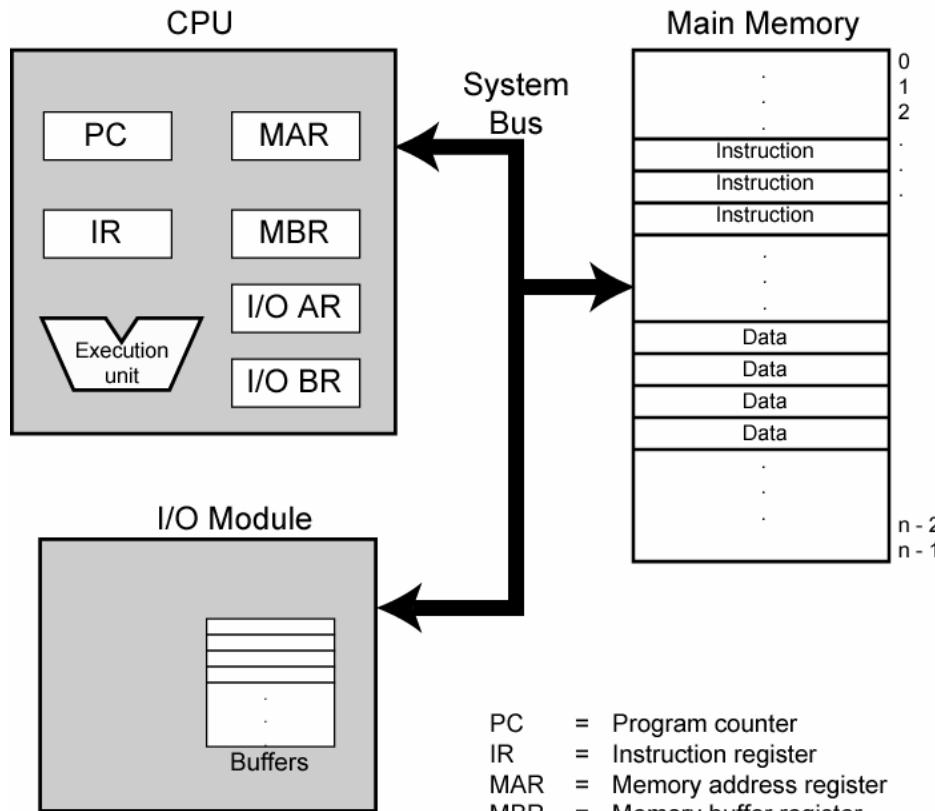
Computer Architecture (Very Fast Recap)

- **What are the major components of a computer?**
- CPU.
- Memory: main and disk.
- I/O components.
- All these are interconnected.
- **What is the main objective of a computer?**
- To execute user programs.
- CPU: executes the instructions.
- Memory: stores data and programs.
- I/O modules: move data between computer and external environment.





Computer Components



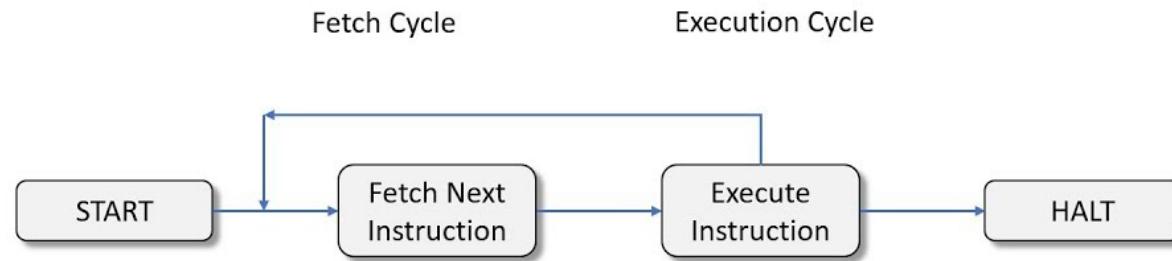
PC = Program counter
IR = Instruction register
MAR = Memory address register
MBR = Memory buffer register
I/O AR = Input/output address register
I/O BR = Input/output buffer register





Instruction Cycle

Computer Instruction Cycle





Instruction Components



(a) Instruction format



(b) Integer format

Program counter (PC) = Address of instruction
Instruction register (IR) = Instruction being executed
Accumulator (AC) = Temporary storage

(c) Internal CPU registers

0001 = Load AC from memory
0010 = Store AC to memory
0101 = Add to AC from memory

(d) Partial list of opcodes





Example of Program Execution

- Goal: add contents of memory word at address 940 to the contents of memory word at address 941, and store the results in the latter location.
- Three instructions (with three fetch and execute stages are needed):
- PC contains 300, address of 1st instruction. This instruction (1940) is loaded into the IR, and PC is incremented, to 301.
- 1st four bits of instruction in IR indicate AC to be loaded from memory. Remaining 12 bits specify memory address = 940. This will load contents of 940 (= 3) into AC.
- Next instruction (= 5941) is fetched from 301, and put in IR. Also, PC is incremented, to 302.
- 1st four bits of instruction in IR indicates add to AC from memory address 941 (which contains 2). We do $3 + 2 = 5$ and put 5 back in AC.
- Next instruction (=2941) is fetched from 302, and put in IR. Also, PC is incremented, to 303.
- 1st four bits of IR instructions indicate to store AC contents to memory 941.
- We are now done.





Program Execution

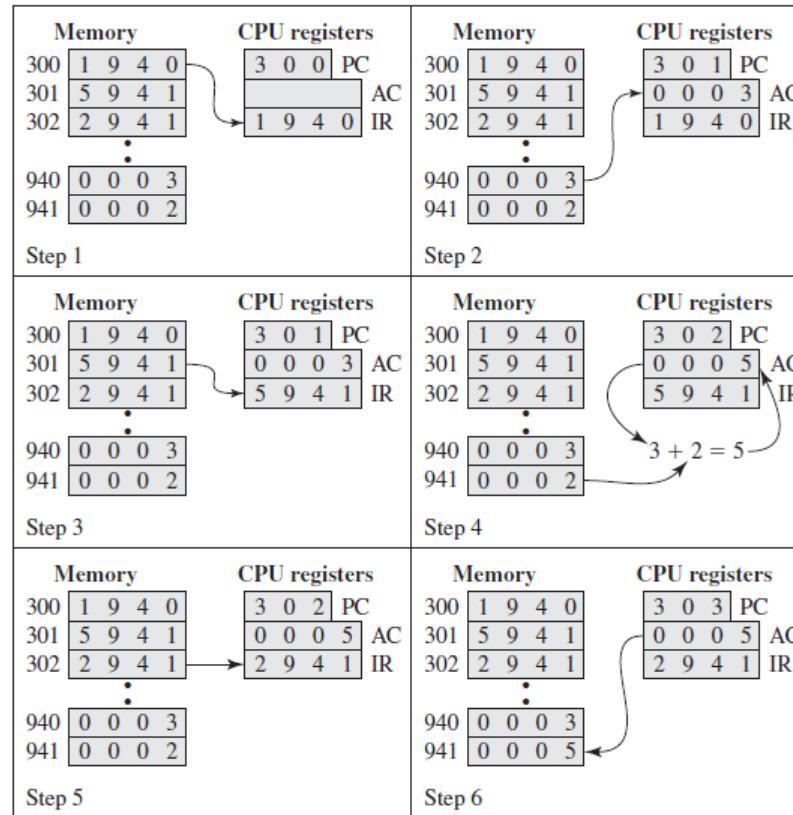


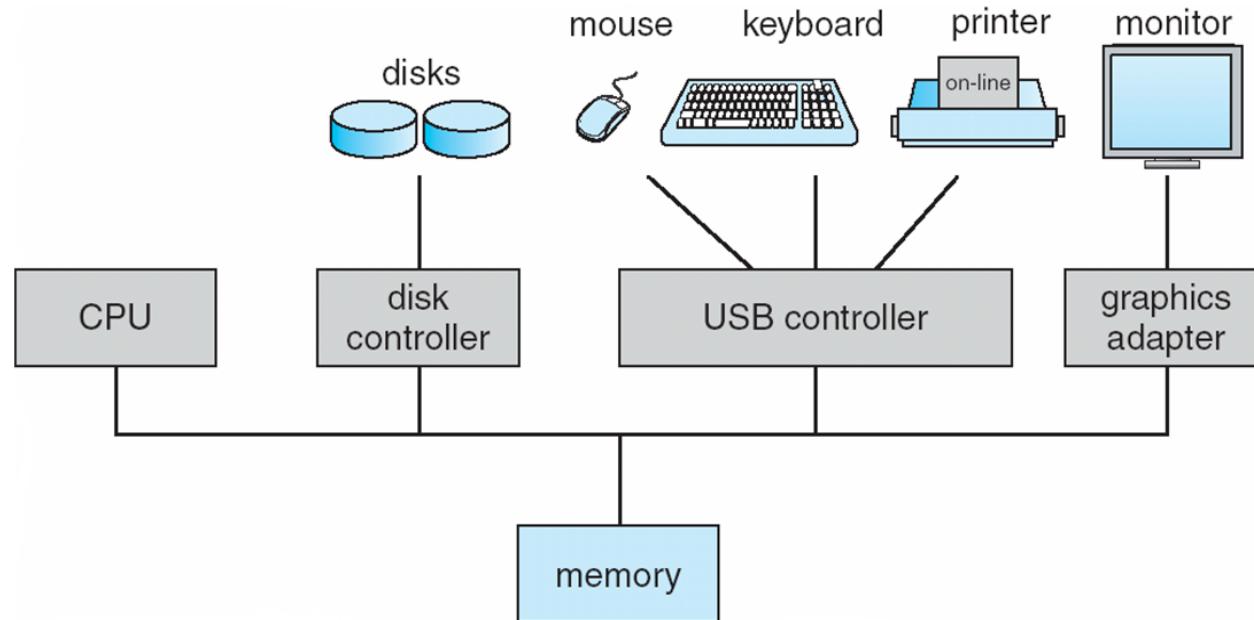
Figure 3.5 Example of Program Execution (contents of memory and registers in hexadecimal)





Computer System Organization

- Computer-system consists of:
 - One or more CPUs, device controllers connected through a common bus providing access to shared memory.
 - Concurrent execution of CPUs and devices competing for memory.





Device Controllers

- Note that each I/O device (disk, USB) has an associated controller.
- **A device controller is a piece of hardware that connects the operating system to the various devices, such as keyboard, mouse, printer etc.**
- Each device controller is in charge of a particular device type.
- **Goal of a device controller?**
- Make I/O data available to the OS.
 - So that the OS can act on it.
- Controller may have its own local buffer to store data.
 - For example, data read by keyboard is put in a buffer, from where it can be read by the OS.
- Controller has a piece of code that accomplishes its task, and it is called a **device driver**.
 - Device driver is software that manages device (hardware) controllers.





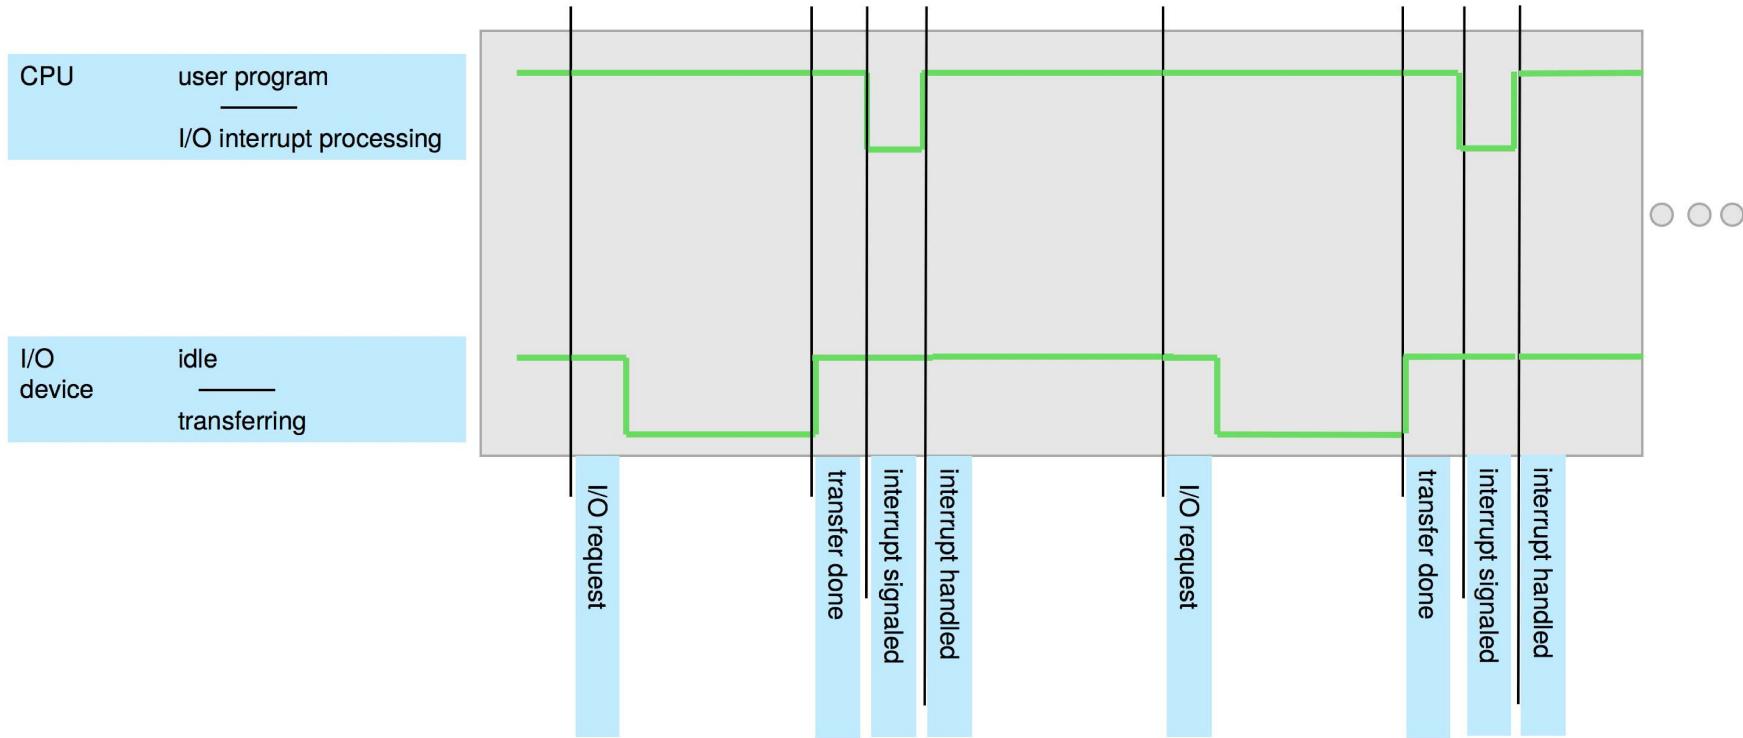
Interrupts

- How does a device driver informs the CPU that it has finished its operation?
- By causing an interrupt.
- **Events in an OS are signaled by interrupts.**
- Interrupts (to an OS) may be signaled by either the hardware or the software (called traps or exceptions or system calls).
- Hardware interrupts CPU by sending a signal.
 - Signal sent through the system bus.
- Software interrupts CPU by executing a system call.
 - This is a piece of code.
- Key concept: **an operating system is interrupt driven.**
 - This is how, for example, the CPU gets to know about a program that wants to start execution.
 - Many other events signaled by interrupts, as we will soon see.





Interrupt Timeline





Interrupt Time Line

- Here's the sequence of events:
 1. CPU is busy executing some instructions of a program.
 2. CPU is interrupted.
 3. It stops it's execution, & transfers to a fixed location in memory, that contains the **interrupt service routine**.
 4. This routine/code is executed.
 5. On completion, CPU resumes execution of interrupted program.
- Table of pointers is maintained.
- Each pointer points to a specific interrupts service routine.
- This table of pointers is stored in low memory locations, & is called an **interrupt vector, say *iv[100]**.
- Vector – “to direct to a desired point”.
- Both Windows and Linux service interrupts in this manner.





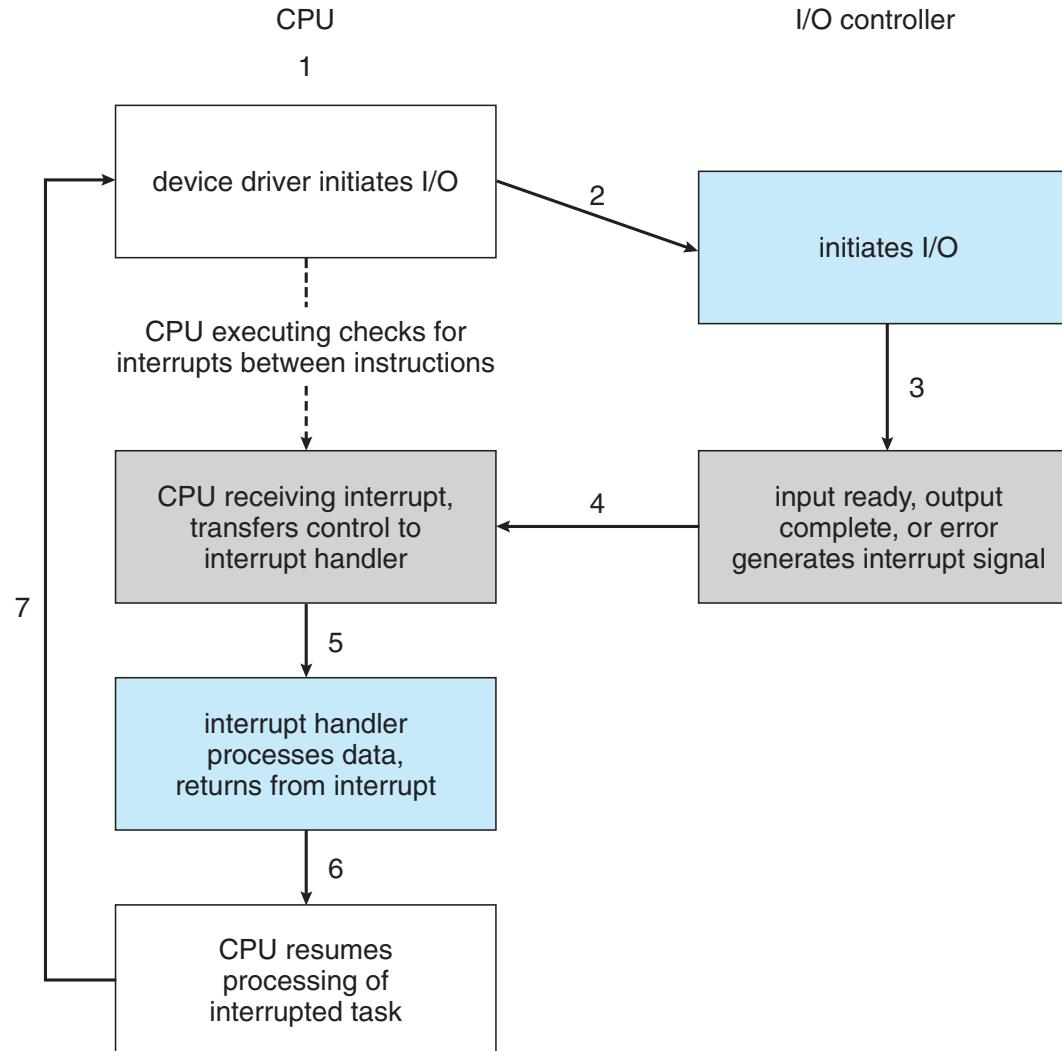
Interrupt Implementation

- CPU hardware has a wire (interrupt request line) that it senses after executing every instruction.
- When CPU detects a device controller has sent a signal on the line, it reads interrupt number and jumps to relevant isr.
- CPU executes the isr code (to service the interrupt).
- **Note that current state of CPU needs to be saved (and restored later).**
- After servicing the interrupt, CPU goes back to what it was doing (execute next instruction of the halted program). **Need a state restore for this.**
- In modern OSes, interrupts are assigned a priority level (which one to service now, which one can wait a bit?).
- Vector number 0: divide error, 6: invalid opcode, 7: device not available, 14: page fault etc.
- Integers here are indexes to array of pointers.





Interrupt I/O Cycle





Types of Interrupts

Program interrupts	As a result of instruction execution e.g. arithmetic overflow, out of bounds instruction.
Timer interrupts	Generated by a timer in the processor to perform functions at regular intervals.
I/O interrupts	Normal I/O completion or error conditions.
Hardware failure interrupts	Power failure.





Storage Structure

- CPU can load instructions only from main memory.
- Any program to be run must be stored there.
- All memory provides an array of bytes.
- Each byte is addressable using a memory location.
- Memory – CPU: load instruction.
- CPU – memory: store instruction.
- Main memory is too small to accommodate all programs.
- Hence, we have external memory, also known as disks.





Storage Structure

- Main memory – only large storage that the CPU can access directly.
 - **Random access** (any memory cell accessed in same amount of time).
 - Typically **volatile** (contents lost on power down).
- Secondary storage – extension of main memory that provides large **nonvolatile** storage capacity. This consists of:
- **Hard disks (HD)** – rigid metal or glass platters covered with magnetic recording material
 - Disk surface is logically divided into **tracks**, which are subdivided into **sectors**
 - The **disk controller** takes care of the interaction between the device and the computer.
- **Solid-state disks (SSD)** – faster than hard disks, nonvolatile.
 - Various technologies.
 - Becoming more popular.
- More on this later in chapters on memory management.

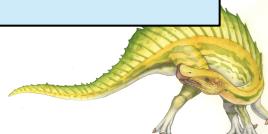




Storage Definitions and Notation Review

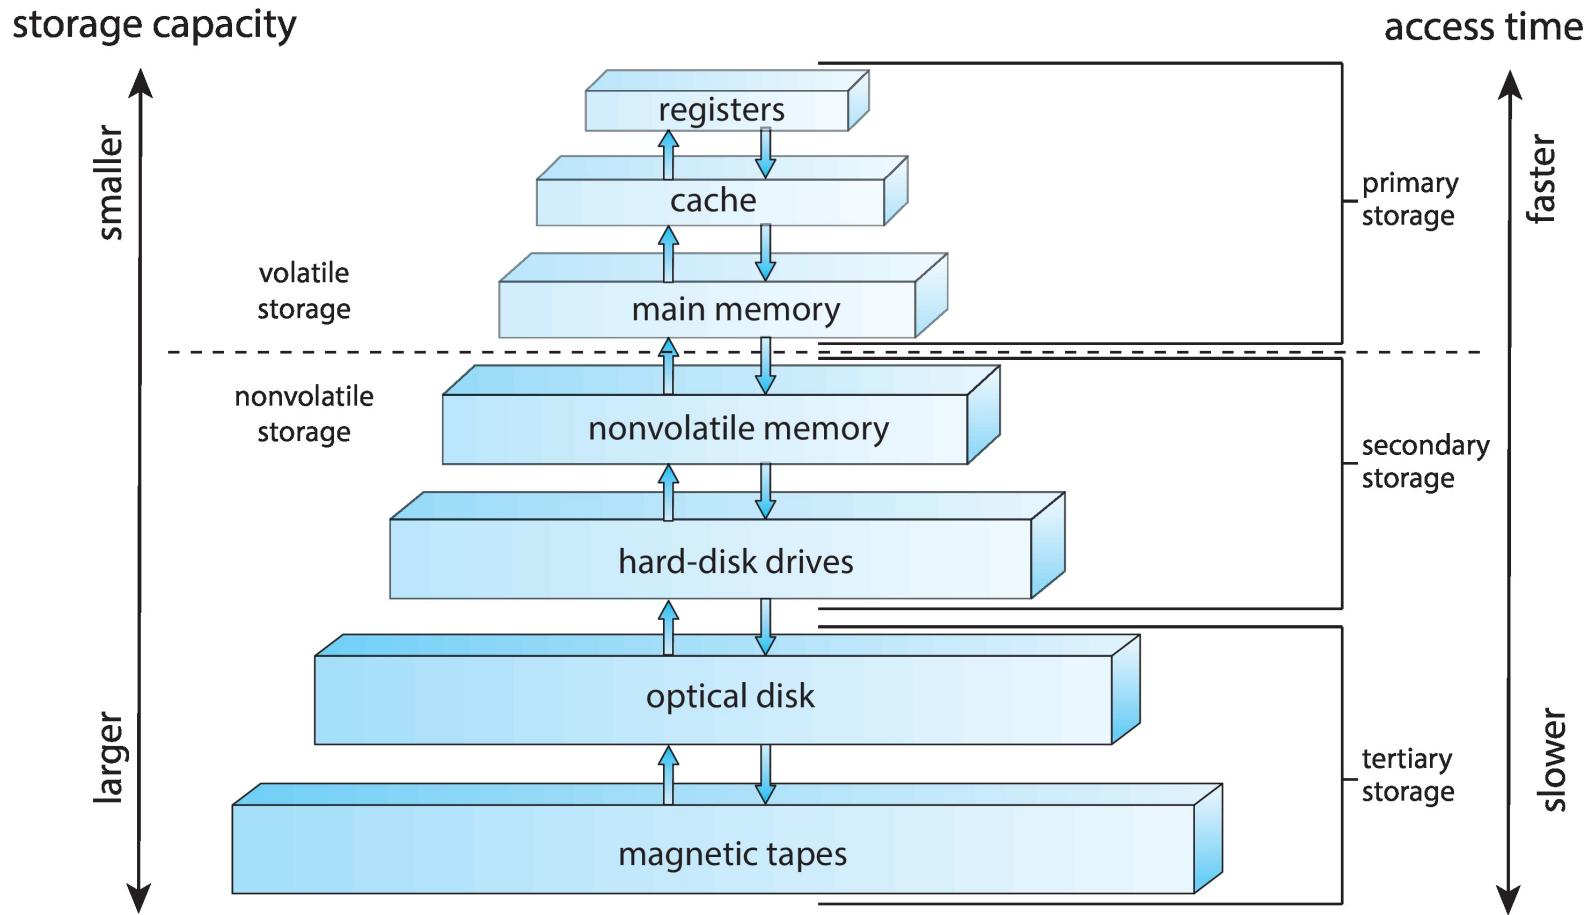
The basic unit of computer storage is the **bit**. A bit can contain one of two values, 0 and 1. All other storage in a computer is based on collections of bits. Given enough bits, it is amazing how many things a computer can represent: numbers, letters, images, movies, sounds, documents, and programs, to name a few. A **byte** is 8 bits, and on most computers it is the smallest convenient chunk of storage. For example, most computers don't have an instruction to move a bit but do have one to move a byte. A less common term is **word**, which is a given computer architecture's native unit of data. A word is made up of one or more bytes. For example, a computer that has 64-bit registers and 64-bit memory addressing typically has 64-bit (8-byte) words. A computer executes many operations in its native word size rather than a byte at a time.

Computer storage, along with most computer throughput, is generally measured and manipulated in bytes and collections of bytes. A **kilobyte**, or KB, is 1,024 bytes; a **megabyte**, or MB, is $1,024^2$ bytes; a **gigabyte**, or GB, is $1,024^3$ bytes; a **terabyte**, or TB, is $1,024^4$ bytes; and a **petabyte**, or PB, is $1,024^5$ bytes. Computer manufacturers often round off these numbers and say that a megabyte is 1 million bytes and a gigabyte is 1 billion bytes. Networking measurements are an exception to this general rule; they are given in bits (because networks move data a bit at a time).





Storage-Device Hierarchy





Caching

- Would have studied this in CSL 211.
- **What is caching?**
- Important principle, performed at many levels in a computer (in hardware, operating system, other software).
- Data in use copied from slower to faster storage temporarily.
- Faster storage (cache) checked first to determine if data is there.
 - If it is, data used directly from the cache (fast).
 - If not, data may be copied to cache and used.
- Challenge: cache smaller than data being cached.
 - Cache management important design problem.
 - Cache size and replacement policy.





Operating-System Operations

- After review of architecture, let's focus back on operating systems.
- Bootstrap program – simple code to initialize the system, load the kernel (discussed earlier).
- No user programs to run, no I/O to process, kernel simply waits for an event, signaled by an interrupt.
- **Interrupt driven** (hardware and software)
 - Hardware interrupt by one of the devices.
 - Software interrupt (**exception** or **trap**):
 - ▶ Software error (e.g., division by zero).
 - ▶ Request for operating system service.
 - ▶ Other process problems -> infinite loop, processes modifying each other or the operating system.





Multiprogramming and Multitasking

■ Multiprogramming.

- Single user cannot keep CPU and I/O devices busy at all times.
- Multiprogramming organizes jobs (code and data) so CPU always has one to execute.
- A subset of total jobs in system is kept in memory.
- One job selected and run via **job scheduling**.
- When it has to wait (for I/O for example), OS switches to another job.
- Advantage: > CPU utilization, keeps users satisfied.

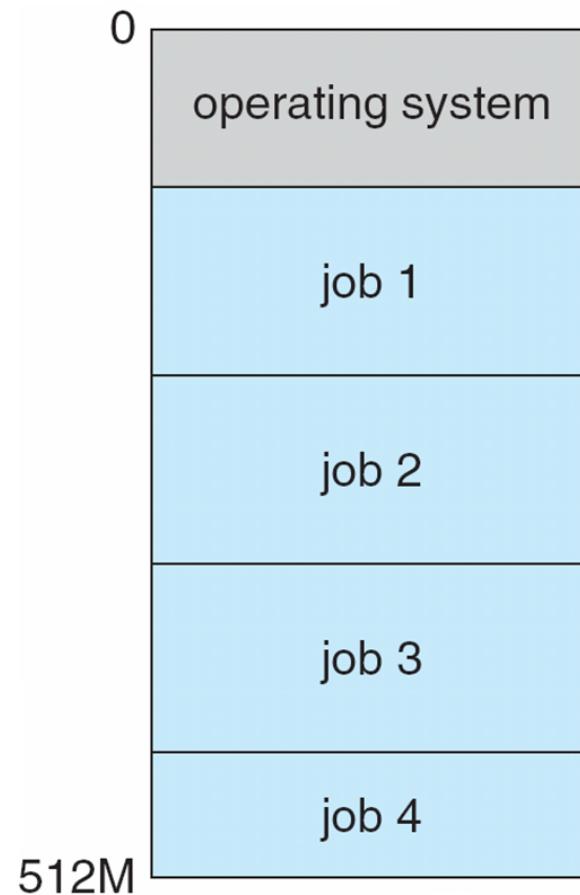
■ Timesharing (multitasking) is logical extension in which CPU switches jobs so frequently that users can interact with each job while it is running, creating **interactive** computing.

- Illusion is that system executes multiple programs “**simultaneously**”.
- **Response time** should be less, say, < 1 second.
- Each user has at least one program executing in memory \Rightarrow **process**.
- If several jobs ready to run at the same time \Rightarrow **CPU scheduling**.
- If processes don't fit in memory, **swapping** moves them in and out to run.
- **Virtual memory** allows execution of processes not completely in memory.





Memory Layout for Multiprogrammed System





Dual Mode

- Software interrupt from a user program that an OS service be provided. This interrupt is called a **system call**.
- The OS and user programs share computer hardware.
- Incorrect or malicious program should not cause other programs, or the OS to operate incorrectly.
 - E.g. process hogging an IO device or the CPU.
- Need a security mechanism that distinguishes between user code and OS code.
- Such a mechanism is called – dual mode operation.
- **Dual-mode** operation allows OS to protect itself and other programs from each other.
 - **Two modes:** **User mode** and **kernel mode**.
 - In kernel mode, CPU can execute all instructions & access all areas of memory (unrestricted access).
 - In user mode, CPU is restricted to certain instructions & accesses certain areas of memory.





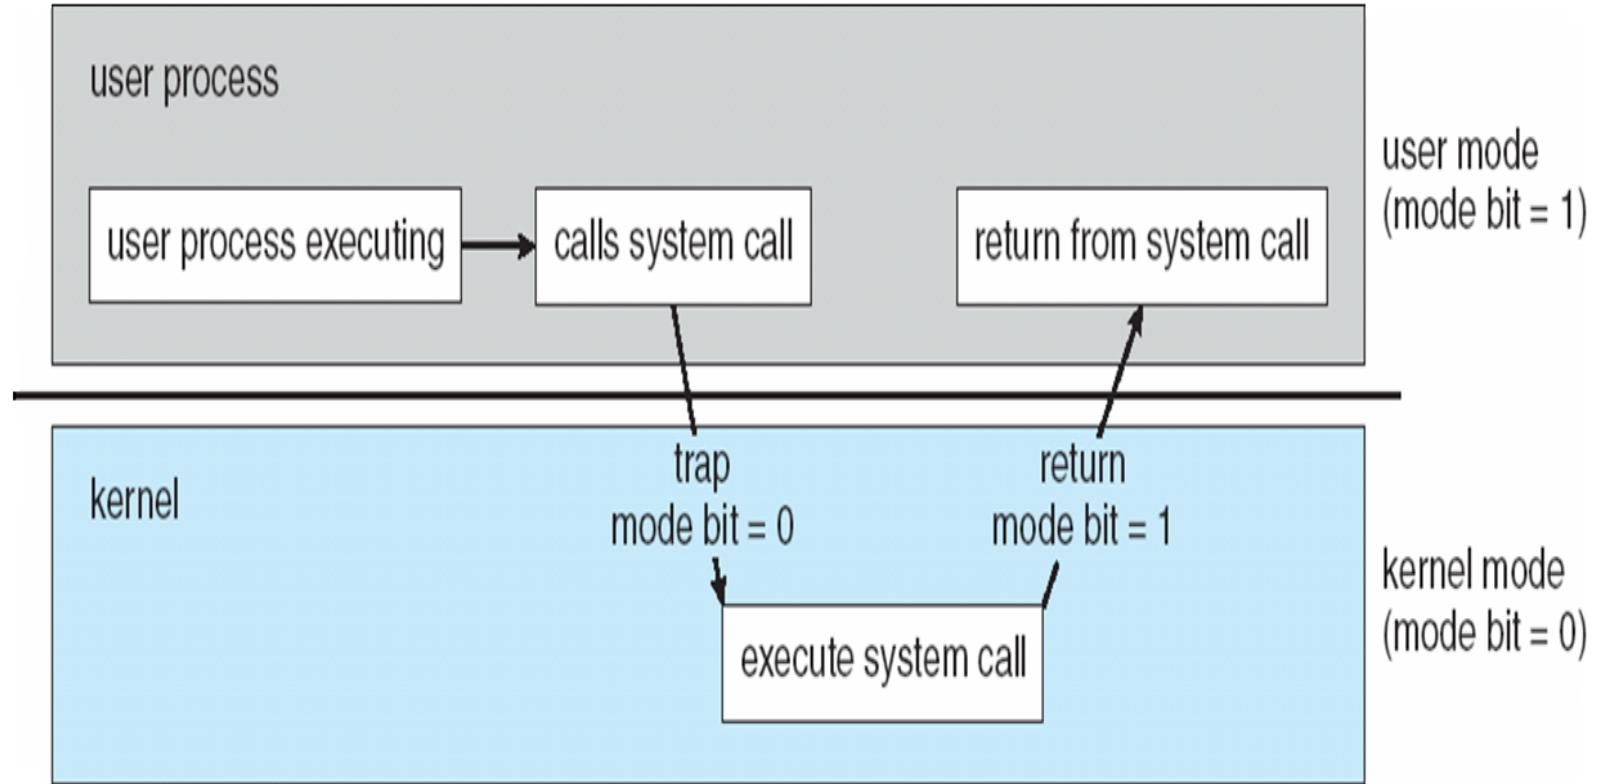
Dual Mode Operation

- **Mode bit** provided by hardware.
 - ▶ Provides ability to distinguish when system is running user code or kernel code.
 - ▶ Some instructions designated as **privileged**, only executable in kernel mode.
 - ▶ System call changes mode to kernel, return from call resets it to user.
- **What if user code tries to execute a privileged instruction?**
- Hardware traps to the OS, which may terminate the process, dump its code to memory for examination (core dump).
- Example of privileged instructions:
 - Go from user → kernel mode.
 - Interrupt management.
 - I/O device management.





Transition from User to Kernel Mode





Source: Julia Evans, Twitter.

drawings.jvns.ca

User space vs. kernel space

JULIA EVANS
@b0rk

the Linux kernel has millions of lines of code

- ★ read + write files
- ★ decide which programs get to use the CPU
- ★ make the keyboard work

When Linux kernel code runs, that's called kernel space

When your program runs, that's user space

time to write a file
that's MY JOB!
YOUR PROGRAM
KERNEL
time for a context switch to kernel space

your program switches back and forth

```
str = "my string"  
x = x+2  
file.write(str) ← ★ switch to kernel space ★  
y = x+4  
str = str*y ← ★ and we're back to user space! ★
```

timing your process

```
$ time find /home  
0.15 user 0.73 system  
↑  
time spent in your process  
↑  
time spent by the kernel doing work for your process
```





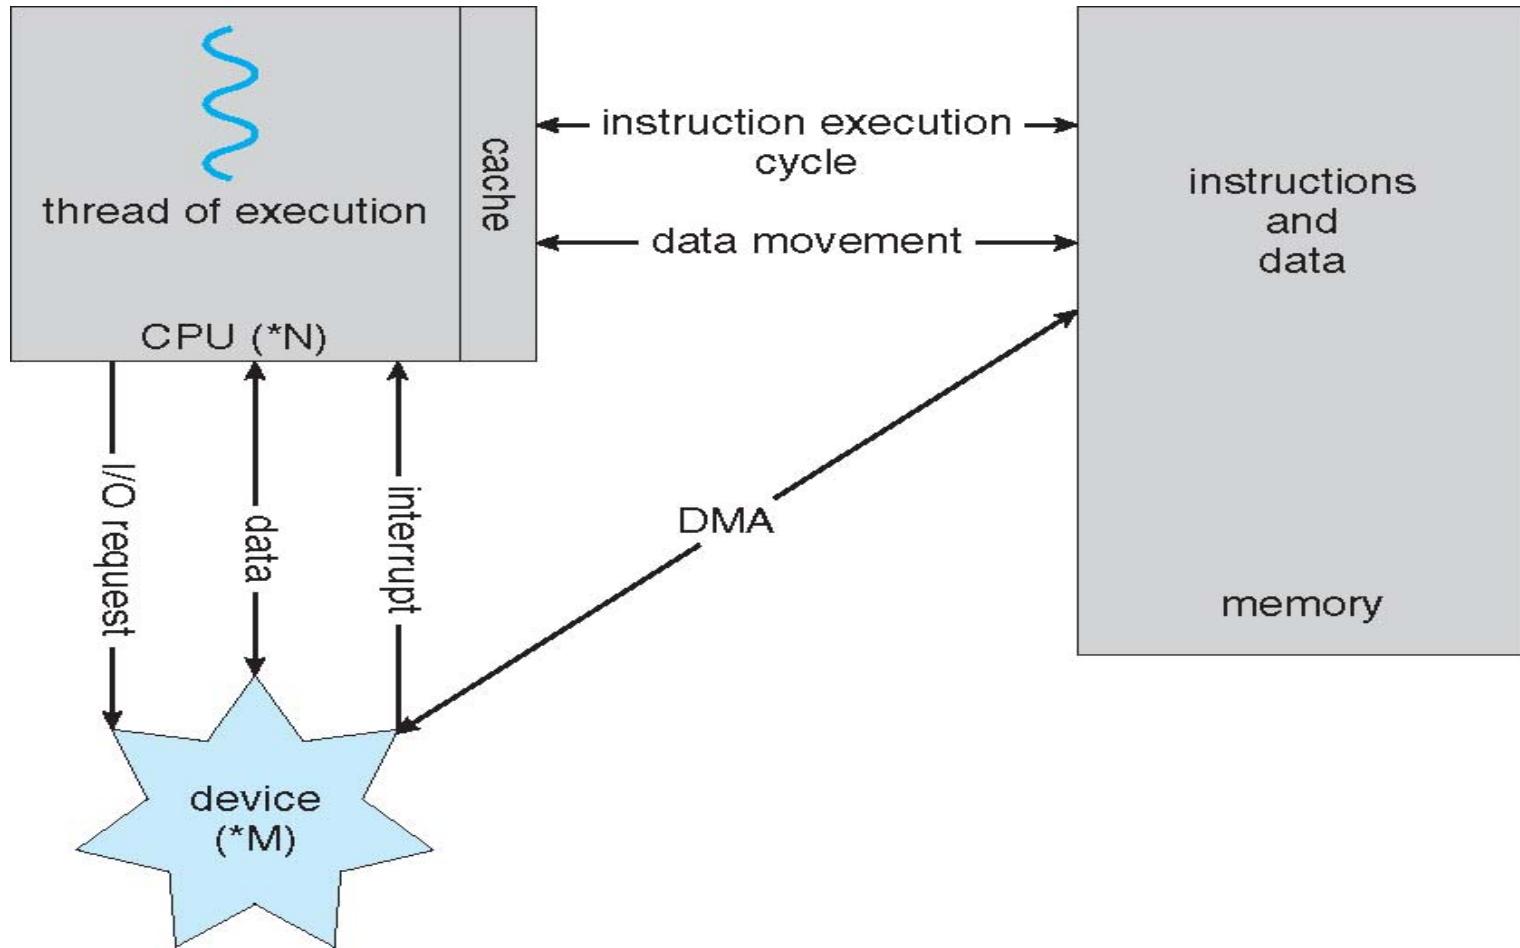
UNIX time utility

- Also try: `time find /usr`.
- Three kinds of time shown in the output:
- User time – time spent executing user code.
- Sys time – time spent executing kernel code.
- Real time – total time spent, including wait times.
- User + Sys = CPU time used by process, across all CPUs.





How a Modern Computer Works



A von Neumann architecture





1. Single CPU systems

- We now briefly discuss some popular computer architectures.
- When hardware was expensive, most systems used to have a single **general-purpose** processor (CPU).
- As opposed to general purpose CPUs, we now have special purpose CPUs.
- For example, for graphics (GPUs) or for capturing motion (motion co-processor) on the iPhone.
 - The motion co-processor collects data from motion sensors (accelerometer, compass), used by various motion based apps.
- In a single CPU machine, at any point of time, only one program may run (instructions of only one program may be executed).
 - However, we can still simultaneously support multiple programs, even with a single CPU. **How?**
 - Using multiprogramming.





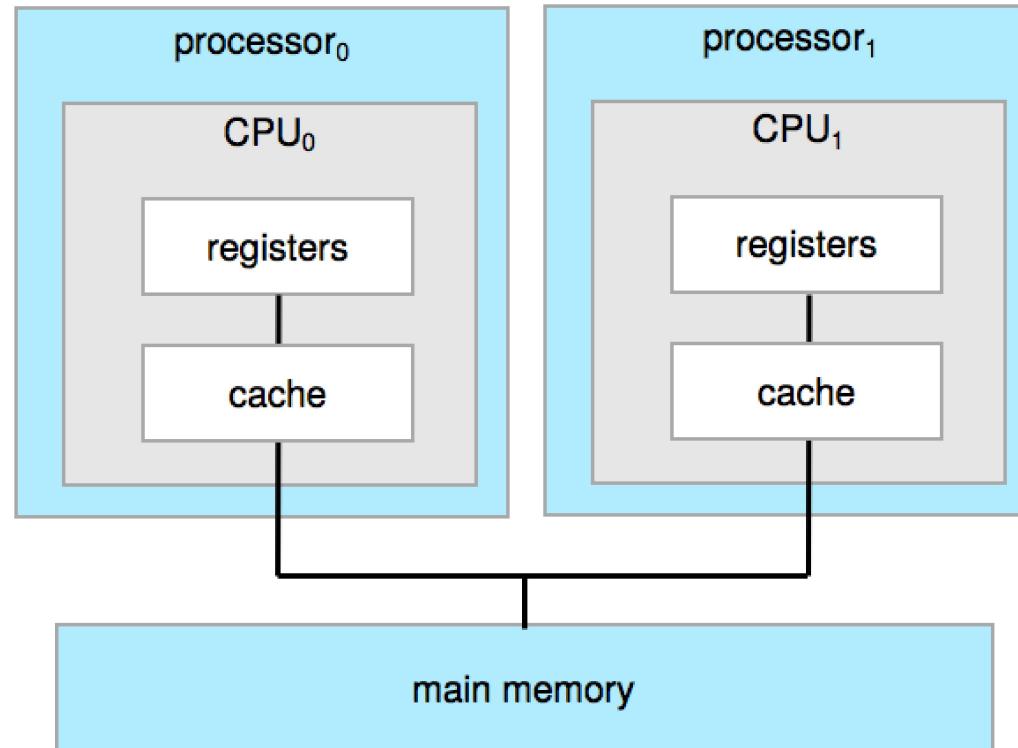
2. Multiprocessor Systems

- **Multiprocessors (parallel)** systems growing in use and importance.
 - Have multiple CPUs on the same machine.
 - **Advantage of this?**
 - Can support more programs at a time.
 - However, the CPUs do share the memory & I/O devices (contention reduces expected gain).
 - Multiprocessors first appeared in servers, but are now available on basic desktops, laptops, & even mobile phones.
- Advantages of multiprocessors:
 1. **Increased throughput** (more work done/unit time).
 2. **Economy of scale** (multiple CPUs sharing memory & IO devices vs. single CPU machines, each with their own memory & IO devices, good for programs that share data).
 3. **Increased reliability** – graceful degradation or fault tolerance. If one of the CPUs breaks down, load of that CPU can be redistributed to remaining CPUs.





Symmetric Multiprocessing Architecture

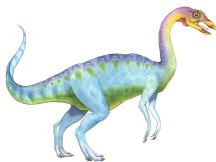




3. Clustered Systems

- Like multiprocessor systems, but multiple (possibly multiprocessor) systems working together, & are generally distributed in nature.
 - A bunch of machines, connected geographically apart, connected by a network.
 - Network may be in a building, in a city, or across countries & continents.
 - Example applications: Facebook, Google.
 - Usually sharing storage via a **storage-area network (SAN)**.
 - Provides a **high-availability** service which survives failures.
 - One master node may monitor many ordinary nodes, that do computation.
 - If one node fails (hardware failure), master node sends request that was running on failed machine to another machine.
 - Request is re-started, only a brief interruption of service for users.





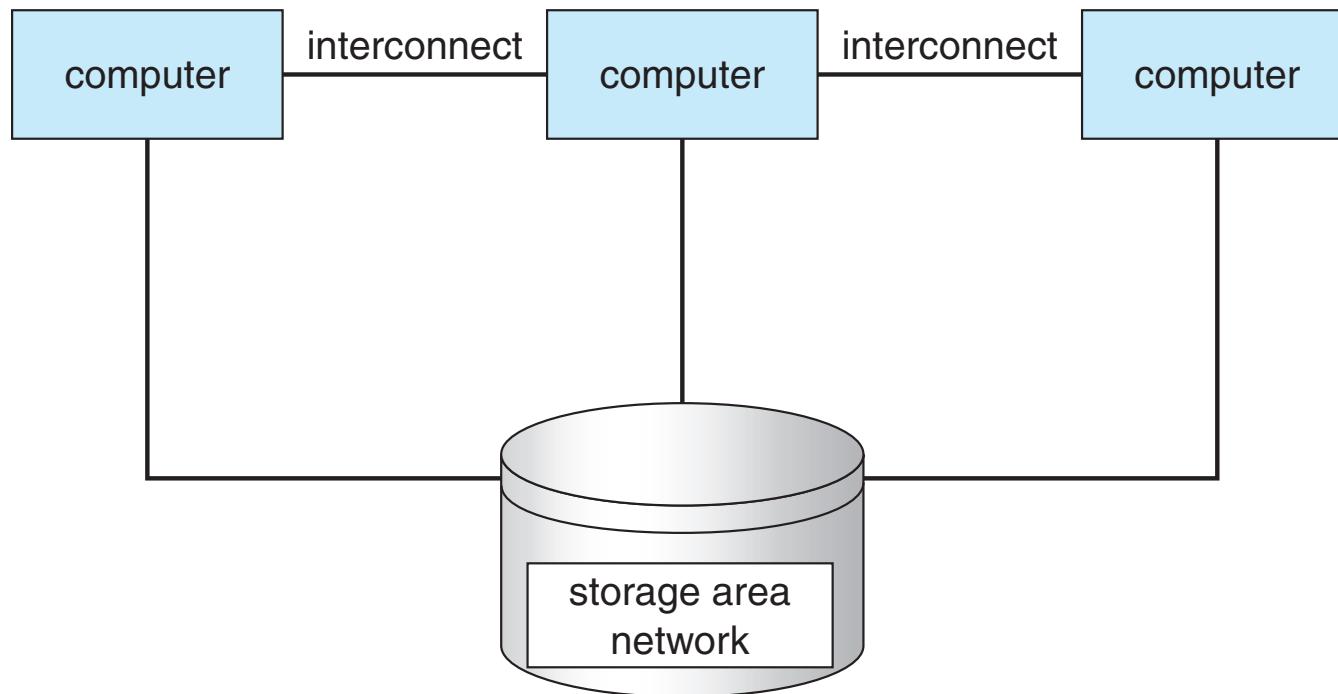
Clusters (continued)

- Clusters generally support **high-performance computing (HPC)**.
- Can use power of hundreds of nodes.
- Can run computationally expensive applications on them, say DNA encoding, weather simulations.
- Divide the application into parts, & run each part in parallel on a separate node/machine (multi threading).
- Applications must be written to use **parallelization**.
- Example application for databases – Oracle Real Application cluster.
- It is a version of Oracle's database that can be run on a parallel cluster.
- Multiple nodes can simultaneously access a database, speeding up operations & providing reliability.
- Storage, however, is shared. Can have a Storage Area Network (SAN).
 - ▶ Distributed storage to which many node can attach.
 - ▶ Cluster software can assign app on SAN to any connected node for running.
- Need to be careful with writes, reads are OK (Distributed Lock Manager).





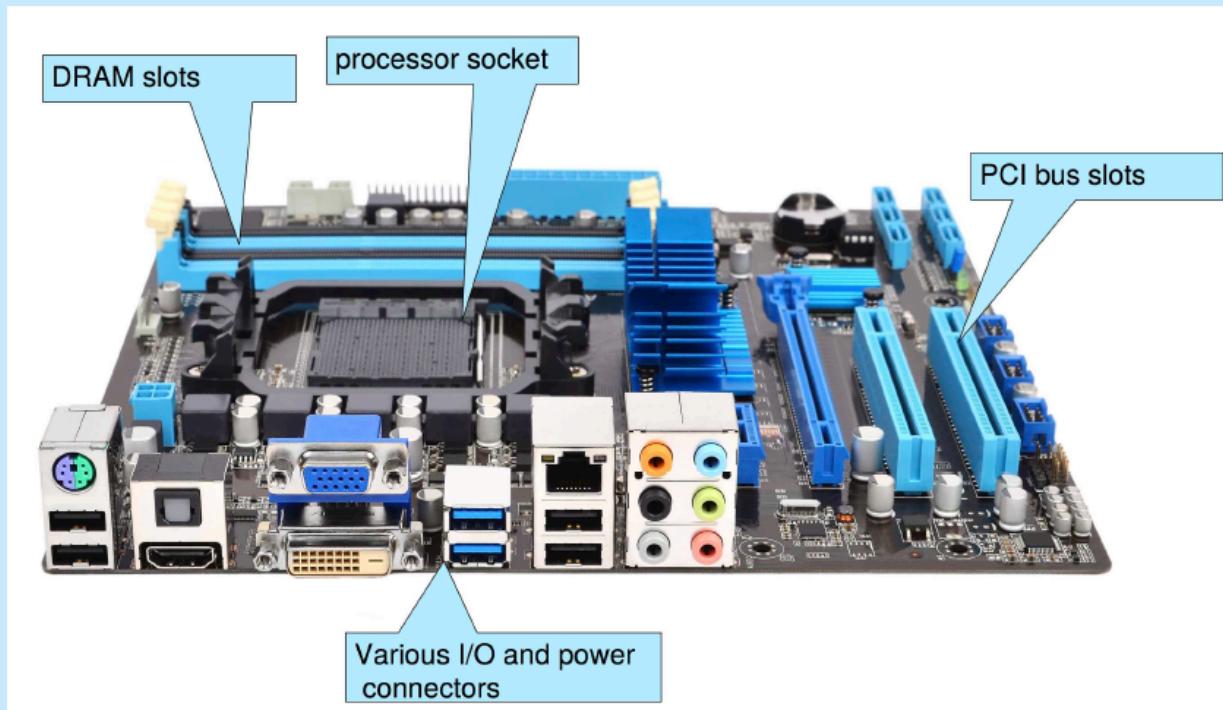
Clustered Systems





PC Motherboard

Consider the desktop PC motherboard with a processor socket shown below:



This board is a fully-functioning computer, once its slots are populated. It consists of a processor socket containing a CPU, DRAM sockets, PCIe bus slots, and I/O connectors of various types. Even the lowest-cost general-purpose CPU contains multiple cores. Some motherboards contain multiple processor sockets. More advanced computers allow more than one system board, creating NUMA systems.





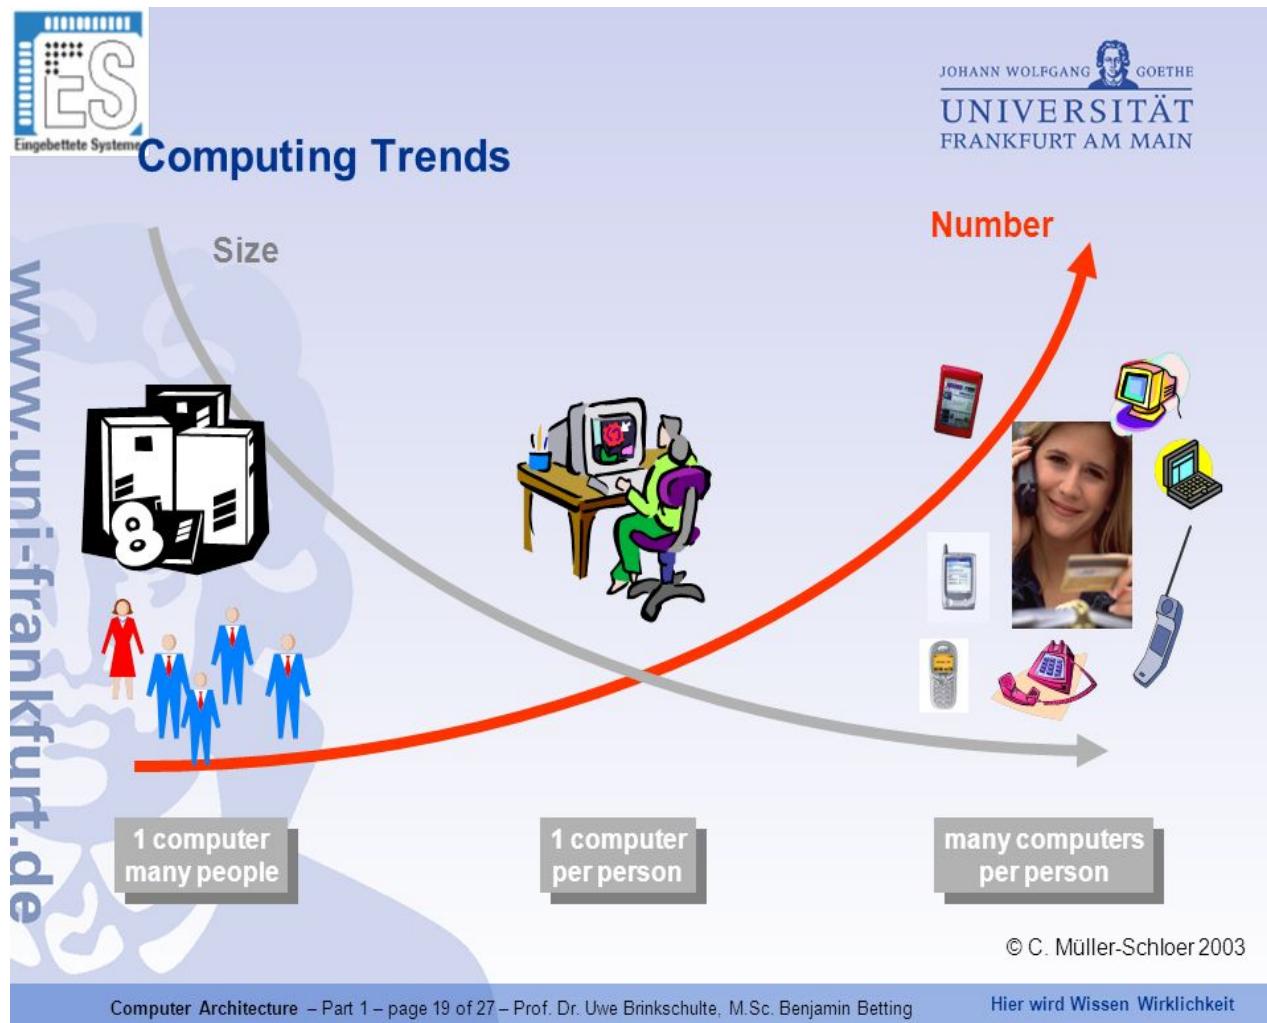
Computing Environments - Traditional

- Stand-alone, general purpose machines.
- For example, a desktop or a laptop at home or in the office.
- Definition is blurred as most systems interconnect with others (networked).
- Earlier, hardware was expensive, so we had time shared systems.
 - One machine accessed by multiple users.
 - **Use multiprogramming to provide fairness among users.**
- Today, time shared systems are uncommon in a home or in one office.
- Think about the quantity -> number of people/machine. **How has it changed over the years?**
- **Multiprogramming can still be used among various programs of a single user (multiple processes for single user).**
- **User processes & system processes each get a slice of the CPU time.**





Source: Computer Architecture, Univ. of Frankfurt, 2003.





Computing Environments - Mobile

- How many mobile phone users in India?
 - 800 million (~ 80 crores) !
- Mobile devices: handheld smartphones, tablets, etc.
- What is the functional difference between them and a “traditional” desktop/laptop?
- More portable, lightweight.
- Less hardware capability (relatively).
- User interface.
- Extra features – more OS features (GPS, accelerometer.)
- Use IEEE 802.11 wireless standard, or cellular data networks for connectivity (EDGE, GPRS, 2G, 3G, 4G, 5G).
- Leaders are: **Apple iOS** and **Google Android**.
- **More on these in chapter 2.**





Computing Environments – Distributed

- Distributed computing.
 - Collection of cluster of separate, possibly heterogeneous, systems networked together.
 - ▶ **Network** is a communications path, **TCP/IP** most common.
 - **Local Area Network (LAN)**: within a room/building/campus.
 - **Wide Area Network (WAN)**: between buildings/cities/countries.
 - **Metropolitan Area Network (MAN)**: link buildings within a city.
 - **Personal Area Network (PAN)**: within a room (between a headset & phone, between wireless mouse & keyboard)
 - Operating System needs to support various networking functions & needs to interface with various network hardware devices.





Distributed Computing - II

- Advantage of a distributed system?
- Because we can (fast internet speeds, up to GBPS).
 - Use the computation power of hundreds & thousands of machines to run large applications.
- Data availability/reliability.

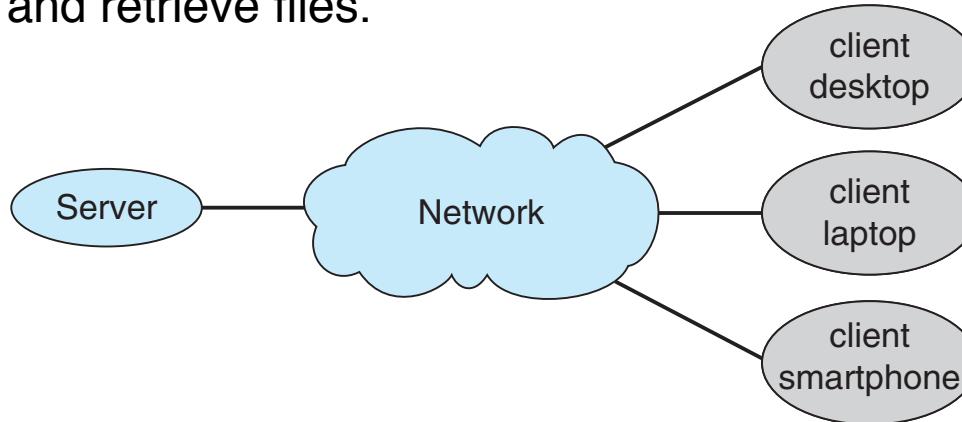




Computing Environments – Client-Server

■ Client-Server Computing.

- Dumb terminals replaced by smart PCs, replaced by dumb terminals again?
- Google server, Facebook server.
- Many systems now use **servers**, responding to requests generated by **clients**.
 - ▶ **Compute-server system** provides an interface to client to request services (i.e., database search).
 - ▶ **File-server system** provides interface for clients to store and retrieve files.





Computing Environments - Peer-to-Peer

- Another model of a distributed system.
- P2P does not distinguish clients and servers.
 - Instead, all nodes are considered peers.
 - Each node acts as client, server or both.
 - Node must join P2P network.
 - ▶ Registers its service with central lookup service on network, or
 - ▶ Broadcast request for service and respond to requests for service
- Examples include Napster and Gnutella, **Voice over IP (VoIP)** such as Skype, torrents.

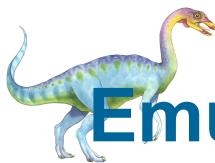




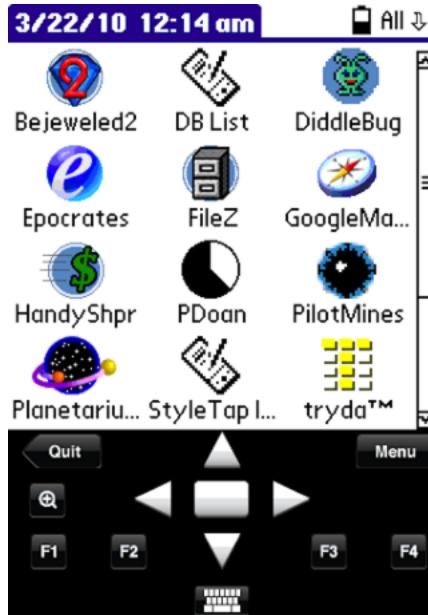
Computing Environments - Virtualization

- Allows whole operating systems to run as applications within other operating systems.
 - Vast and growing industry.
- **Is a form of emulation**, used when source CPU type different from target CPU type.
 - Some time back, Apple changed CPU architecture, from PowerPC to Intel x86.
 - Included a utility called “Rosetta” to let users run PowerPC programs on Intel.
 - Generally, this method is slow.
 - Every instruction needs to be translated from target -> source architecture.
 - Examples: Nintendo emulators (Android), Palm emulator (Android), DOS Box (Linux, Mac OS).





Emulators for Nintendo and Palm on Android





Virtualization - II

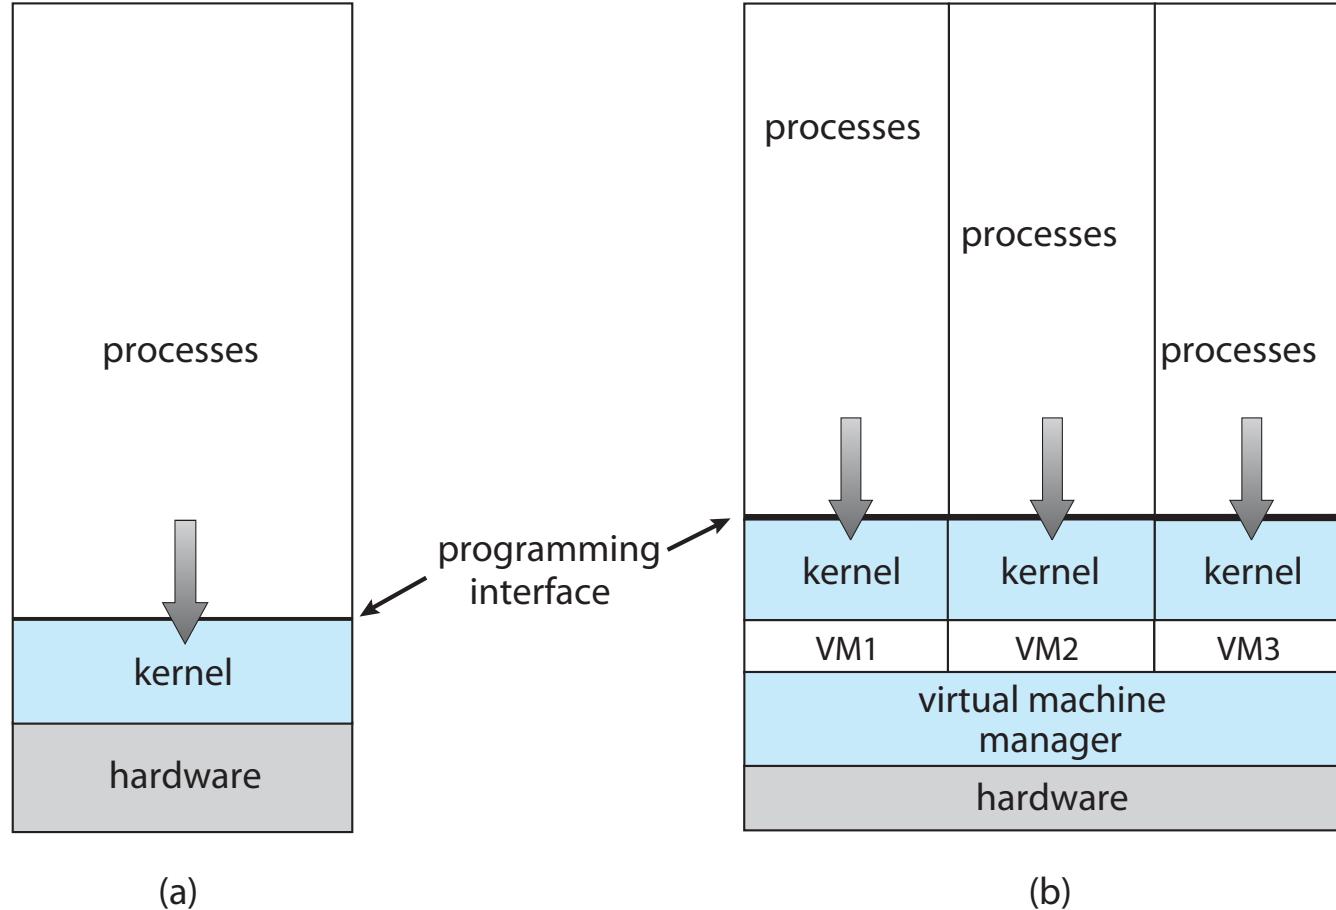
■ **Virtualization** – OS natively compiled for CPU, running **guest** Oses, which may also be natively compiled.

- Same set of hardware shared among multiple OSes.
- To enable a computing environment to run multiple independent systems at the same time.
- Both OSes need same hardware to run.
- Consider VMware running WinXP guests, each running applications, all on native WinXP **host** OS: many users using same system.
- This was how virtualization started initially, at IBM in the 1960s and 1970s.
- **VMM** (virtual machine Manager) software provides virtualization services.
- It has an application that runs on the host OS.
- Over it, many copies of guest OSes can be run.
- Next slide shows a picture.





Computing Environments - Virtualization





Computing Environments - Virtualization

- Use cases involve laptops and desktops running multiple OSes for exploration.
 - Apple laptop running Mac OS X host, Windows as a guest.
 - Developing apps for multiple OSes without having multiple systems.
 - Quality Assurance testing applications without having multiple systems.





Computing Environments – Cloud Computing

- **What is cloud computing?**
- **Delivers computing, storage, even apps as a service across a network.**
- **Any example of cloud based applications that you can think of?**
- Logical extension of virtualization because it uses virtualization as the base for its functionality.
 - Example: Amazon **Elastic Compute Cloud 2 (EC2)** has thousands of servers, millions of virtual machines, petabytes of storage available across the Internet, pay based on usage.
 - Another example is: Microsoft Azure.
- Many types of clouds in the market:
 - **Public cloud** – available via Internet to anyone willing to pay.
 - **Private cloud** – run by a company for the company's own use.
 - **Hybrid cloud** – includes both public and private cloud components.
 - Software as a Service (**SaaS**) – one or more applications available via the Internet (i.e., word processor).
 - Infrastructure as a Service (**IaaS**) – servers or storage available over Internet (i.e., storage available for backup use).
- Advantage of cloud computing? Disadvantage?





Computing Environments – Real-Time Embedded Systems

- Embedded systems are very common today.
- It is a “computer system with a dedicated/specialized function within a larger mechanical or electrical system” – source: Wikipedia.
- Examples: microwave ovens, music players, parts of automobiles, sensors & so on.
- Many other special computing environments as well
 - They have simple OSes, with little or no user interface.
- Usage expanding with the internet (e.g. a sensor network in a smart home). New field -> Internet of Things (IoT)
- Many times, embedded systems run a Real-time OS, that has well-defined fixed time constraints
 - Processing **must** be done within constraint.
 - Correct operation only if timing constraints met.





Real-Time Systems

- Examples:
 - Braking system of an airplane.
 - Sensor network that warns about enemy movement in a military application.
- Need the results of computation in time, else, we may have a disaster.
- In a real-time OS, processes must be completed within a deadline, else, the usefulness is zero.
- Note that in non real-time systems, it is desirable (not essential) to finish tasks early.
- Also note that a real-time is not the same as an “as fast as possible system”.
- Such systems require special scheduling algorithms, as we shall see in a later chapter.





1. Process Management

- What is a process?
- A process is a program in execution. It is a unit of work within the system. Program is a **passive entity**, process is an **active entity**.
- Process needs resources to accomplish its task
 - CPU, memory, I/O, files.
- Process termination requires reclaim of resources.
- Single-threaded process has one **program counter** specifying location of next instruction to execute
 - Process executes instructions sequentially, one at a time, until completion.
- Multi-threaded process has one program counter per thread.
- Typically system has many processes, some user, some operating system running concurrently on one or more CPUs (activity monitor/task manager)
 - Concurrency by sharing the CPUs among the processes / threads.





Process Management Activities

The operating system is responsible for the following activities in connection with process management:

- Creating and deleting processes (user + system).
- Suspending and resuming processes.
- Scheduling processes and threads.
- Providing mechanisms for process synchronization.
- Providing mechanisms for process communication.





2. Memory Management

- Memory management determines what is in memory and when.
 - Goal is: optimizing CPU utilization and computer response to users.
- Memory management activities:
 - Keeping track of which parts of memory are currently being used and by whom.
 - Deciding which processes (or parts thereof) and data to move into and out of memory.
 - Allocating and de-allocating memory space as needed.





3. Storage Management

- OS provides uniform view of information storage.
 - Abstracts physical properties to logical storage unit - **file**.
 - Each storage medium is controlled by device (i.e., disk drive, tape drive)
 - ▶ Varying properties include access speed, capacity, data-transfer rate, access method.
- File-System management.
 - Files usually organized into directories.
 - Access control on most systems to determine who can access what
 - OS activities include
 - ▶ Creating and deleting files and directories.
 - ▶ Disk scheduling
 - ▶ Mounting and unmounting directories.
 - ▶ Storage allocation.
 - ▶ Backup files onto stable (non-volatile) storage media.





IBM Disk from 1956 (business insider)





Performance of Various Levels of Storage

Level	1	2	3	4	5
Name	registers	cache	main memory	solid state disk	magnetic disk
Typical size	< 1 KB	< 16MB	< 64GB	< 1 TB	< 10 TB
Implementation technology	custom memory with multiple ports CMOS	on-chip or off-chip CMOS SRAM	CMOS SRAM	flash memory	magnetic disk
Access time (ns)	0.25 - 0.5	0.5 - 25	80 - 250	25,000 - 50,000	5,000,000
Bandwidth (MB/sec)	20,000 - 100,000	5,000 - 10,000	1,000 - 5,000	500	20 - 150
Managed by	compiler	hardware	operating system	operating system	operating system
Backed by	cache	main memory	disk	disk	disk or tape





4. I/O Subsystem

- One purpose of OS is to hide peculiarities of hardware devices from the user.
 - Provide users with a simple interface for performing I/O.
- I/O subsystem responsible for
 - Memory management of I/O including buffering (storing data temporarily while it is being transferred), caching (storing parts of data in faster storage for performance).
 - Device drivers for specific hardware devices, that provide the above interface.
- Chapter 13 discusses how the I/O system:
 - Interfaces to other parts of the OS.
 - Manages I/O devices.
 - Detects I/O completion.





Open Source Operating Systems

- Two broad categories of operating systems:
 - Windows – closed source.
 - Linux – open source.
- What do we mean by open source?
- They are available in source code format (e.g. C).
- Hence, they can be changed & compiled.
- Can also be copied and distributed for free.
- Closed source, on the other hand, is available in compiled binary code.
- Hard to derive source code from it.
- Apple's OSX has a hybrid approach.
 - Kernel (called Darwin) is open source.
 - Some components are close source.





Open Source Operating Systems

- Advantage of open source operating systems?
- Ideal for learning operating systems.
 - Can go into the detailed code of each module.
- Anyone can make a contribution to the OS to make it better.
 - For example: improve a scheduling algorithm, cache policy.
- They are generally free.
 - For original OS & upgrades as well.
- Dedicated community of programmers, online forums to help out users.
- Arguably more secure.
 - Since source code is in public domain, anyone can suggest changes that enhance the security of the OS.





A Bit of History of Open Source

- In the 1950s, MIT had a Tech Model club, that used to write programs & leave them in drawers, inviting other members to look at the code & suggest changes.
- Companies, such as the Digital Equipment Corporation (DEC) accepted contributions from people for source code, collected such programs, & distributed them to interested users.
- Companies such as Microsoft, Apple did not like this. **Why?**
- They wanted to make money of the OS.
- Initial effort to design the OS.
- Later, only need to make copies of it.
- Some tied the OS to the hardware.
- To use iOS, need to buy an iPhone.
- For Android, there are many phone options.





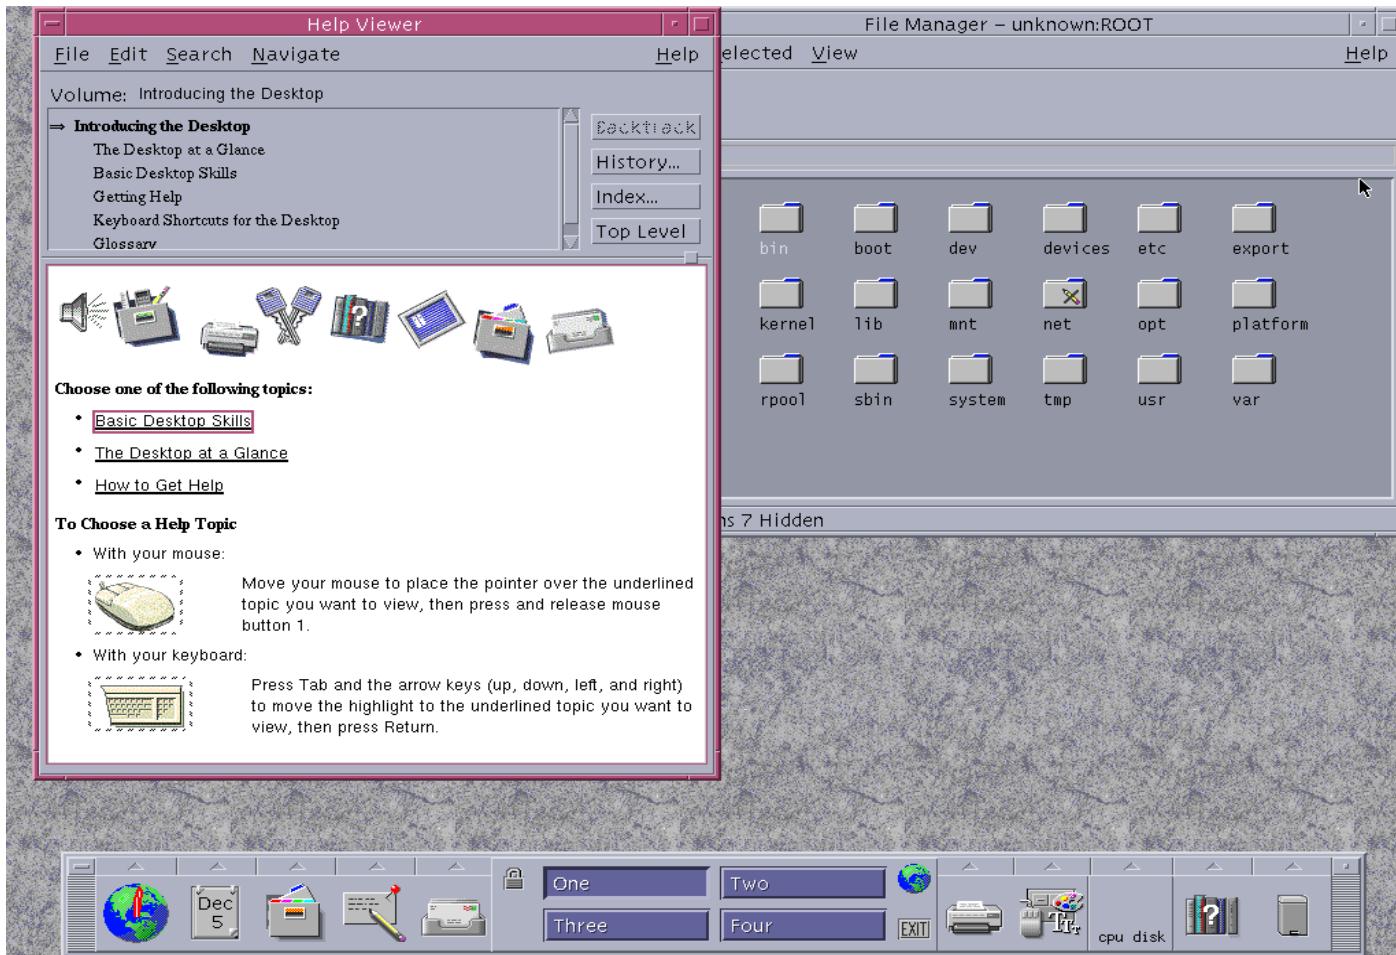
Open Source Operating Systems

- **Linux**, started by Linus Torvalds, a research student from Finland.
- Comes in many flavors – Red Hat, Fedora, Debian, Ubuntu.
- Many Linux projects are provided at: <http://distrowatch.com>.
- Can be run on a Windows/Apple machine using virtualization.
- Install Vmware player or Virtual Box tool.
- Download a Linux distribution.
- Run it as an application.
- **Unix**, on which Apple's OS is based.
- Some open source & some closed source parts.
- Darwin kernel is available for download at: <https://opensource.apple.com>.
- **Solaris**, by Sun Systems.
- This is another Unix variant.





Solaris OS (Wikipedia)





Open Source Software

- Open Office: equivalent of Word, Excel, Power Point.
- Emacs : programming editor.
- Mozilla Firefox: web browser.
- My SQL: database.
- Scilab: equivalent of Matlab, for numerical computation.
- VLC player: music, videos.
- Gimp: image editor.
- Many more.





Open-Source Operating Systems - Recap

- Operating systems made available in source-code format rather than just binary **closed-source**
- Started by **Free Software Foundation (FSF)**, which has “copyleft” **GNU Public License (GPL)**
- Examples include **GNU/Linux** and **BSD UNIX** (including core of **Mac OS X**), and many more
- Can use VMM like VMware Player (Free on Windows), Virtualbox (open source and free on many platforms - <http://www.virtualbox.com>)
 - Use to run guest operating systems for exploration





Summary of Chapter 1

- The OS is a software that – (a) manages the hardware resources and (b) provides a convenient environment for users to run programs.
- Some broad goals of a general purpose OS: convenient user interface, efficient utilization of resources.
- OS has many jobs: process management, memory management, storage management, I/O management, protection & security.
- A (very) brief overview of a computer system and various architectures: single CPU, multiple CPU, distributed, peer-to-peer, cloud, real-time, multimedia, mobile.
- An introduction to some OS features – multiprogramming/multitasking, dual-mode (user & kernel).
- Discussed various architectures: uniprocessor, multiprocessor, distributed.
- Discussed various computing environments: traditional, mobile, cloud.
- Talked about open source operating systems, mostly Unix/Linux.



End of Chapter 1

